

What's New in LLVM

Session 417

Jim Grosbach

Manager, LLVM CPU Compiler Team

Apple LLVM Compiler



Apple LLVM Compiler

LLDB

Swift

CoreImage

JavaScript FTL

OpenGL

Metal

Xcode Static Analyzer

Xcode Objective-C Modernization

Xcode Code Completion



Tools for Modernization and Performance



New in LLVM

64-bit iOS support

Objective-C modernization tool

User-defined modules

Profile-guided optimization

Vectorizer advancements

JavaScript Fourth Tier LLVM

64-Bit iOS Support

Building for 64-Bit iOS

Included in standard architectures: armv7, arm64

Just rebuild with Xcode 6

Can continue to deploy back to iOS 4.3

iOS Simulator fully supports 64-bit development

Static libraries must be built for arm64, including third-party code

Migration Hints

Function type checking

Objective-C BOOL type

Type size independence

Missing Prototypes

All functions must have a prototype for ARM64

```
int my_function(int val) {  
    return other_function(val);  
}
```

Missing Prototypes

All functions must have a prototype for ARM64

```
int my_function(int val) {  
    return other_function(val);  
}
```

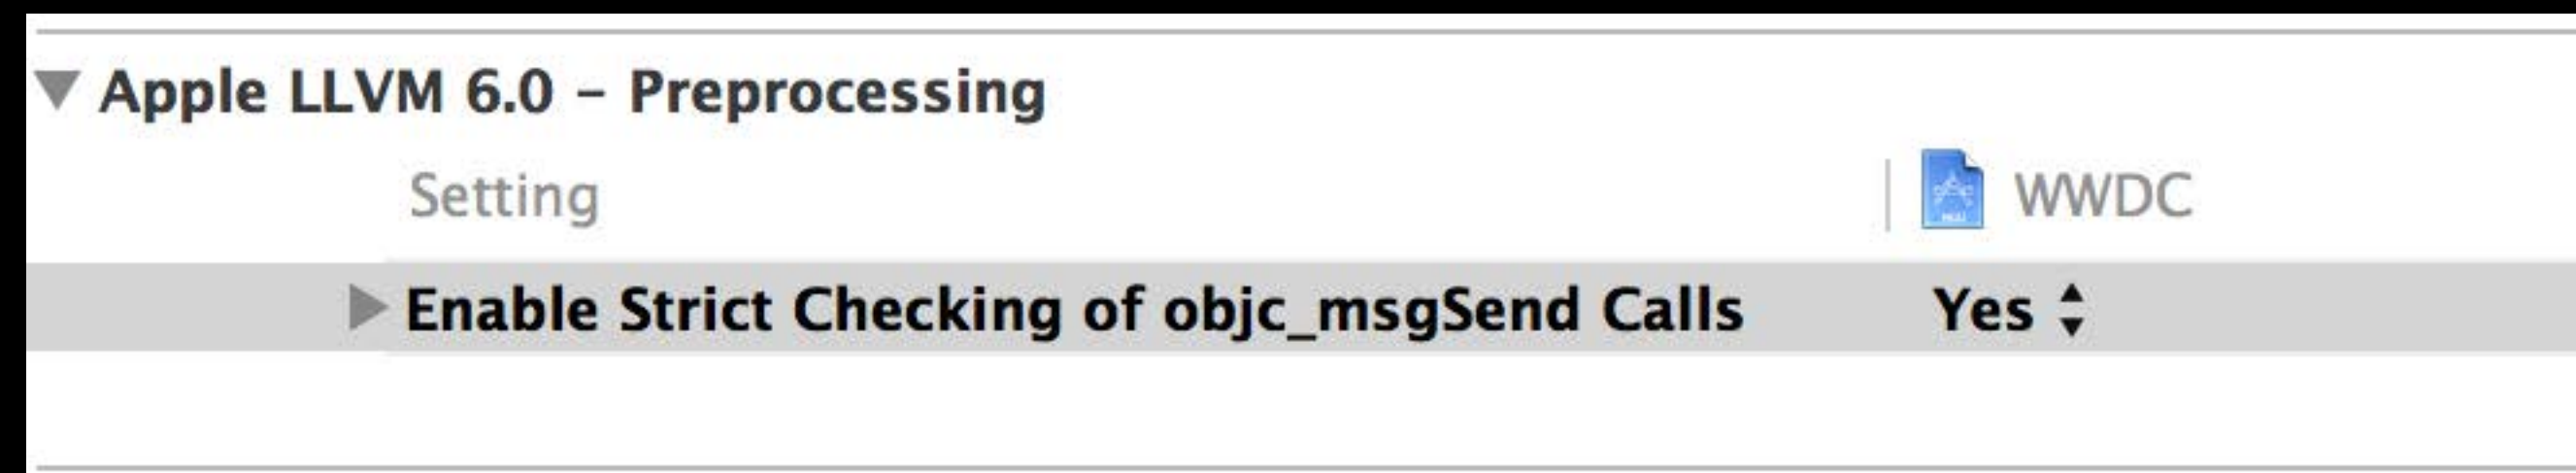
error: implicit declaration of function 'other_function' is invalid in C99
[-Werror,-Wimplicit-function-declaration]
 return other_function(val);
 ^

Strict Checking of objc_msgSend



objc_msgSend without a typecast is usually an error

Recommended build setting for strict checking



objc_msgSend Example

```
#include <objc/message.h>
void foo(void *object) {
    objc_msgSend(object, sel_getUid("foo:"), 5);
}
```

objc_msgSend Example

```
#include <objc/message.h>
void foo(void *object) {
    objc_msgSend(object, sel_getUid("foo:"), 5);
}
```

error: too many arguments to function call, expected 0, have 3
objc_msgSend(object, sel_getUid("foo:"), 5);

~~~~~ ^~~~~~

# objc\_msgSend Example

```
#include <objc/message.h>
void foo(void *object) {
    typedef void (*send_type)(void *, SEL, int);
    send_type func = (send_type)objc_msgSend;
    func(object, sel_getUid("foo:"), 5);
}
```

# Objective-C Boolean Type

`B00L` is a `_Bool` in 64-bit iOS and `signed char` in 32-bit iOS

```
int foo(B00L b) {  
    return b + 1;  
}
```

```
int value = foo(-1);    // 0 in 32-bit code, 2 in 64-bit code
```

# Pointer Casting

Pointer size change can expose latent bugs

`int` is 32 bits, but pointers are 64 bits

```
void *foo(int i) {  
    return (void*)i;  
}
```



# Pointer Casting

Pointer size change can expose latent bugs

`int` is 32 bits, but pointers are 64 bits

```
void *foo(int i) {  
    return (void*)i;  
}
```

warning: cast to 'void \*' from smaller integer type 'int'

[-Wint-to-void-pointer-cast]

```
    return (void*)i;
```

^

# Size Safe Types for Pointers

`intptr_t` and `uintptr_t` for saving pointer values to integers

`size_t` for array indices

`ptrdiff_t` for pointer differences

```
void *foo(intptr_t i) {  
    return (void*)i;  
}
```

# Structure Layouts

Object sizes can change between 32-bit and 64-bit iOS

For example: On-disk formats, network protocols, etc.

```
struct foo {  
    long a; // 4 bytes in 32-bit iOS, 8 bytes in 64-bit iOS.  
    int b;  // 4 bytes always.  
};
```

# Summary: Building for 64-Bit iOS

Enabled by default

Easy to adopt

Compiler assistance to find and resolve issues

# Objective-C Modernization Tool

# Modernizations

Property attribute annotations

Designated initializers

instancetype

Objective-C literals

Objective-C subscripting

Enumeration macros

*Demo*

Xcode Objective-C modernization tool

# Summary: Objective-C Modernization Tool

Many modernizations to the Objective-C language

Xcode modernization tool makes it easy to adopt best practices



# User-Defined Modules

Bob Wilson

Manager, LLVM Core Team

# Modules: Background

Modern alternative to precompiled headers

Precompiled headers speed up compilation, but are not flexible

Textual includes have problems

- Multiple inclusion
- Fragile headers

# Modules: Background

Modern alternative to precompiled headers

Precompiled headers speed up compilation, but are not flexible

Textual includes have problems

- Multiple inclusion
- Fragile headers

```
#define count 100  
#import <Foundation/Foundation.h>
```

# Modules: Background

Modern alternative to precompiled headers

Precompiled headers speed up compilation, but are not flexible

Textual includes have problems

- Multiple inclusion
- Fragile headers

```
#define count 100  
  
. . .  
@interface NSArray : NSObject  
- (NSUInteger)count;  
- (id)objectAtIndex:(NSUInteger)index;  
@end  
  
. . .
```

# Modules: Background

Modern alternative to precompiled headers

Precompiled headers speed up compilation, but are not flexible

Textual includes have problems

- Multiple inclusion
- Fragile headers

```
#define count 100  
  
. . .  
@interface NSArray : NSObject  
- (NSUInteger)100;  
- (id)objectAtIndex:(NSUInteger)index;  
@end  
  
. . .
```

# Not Just for System Frameworks

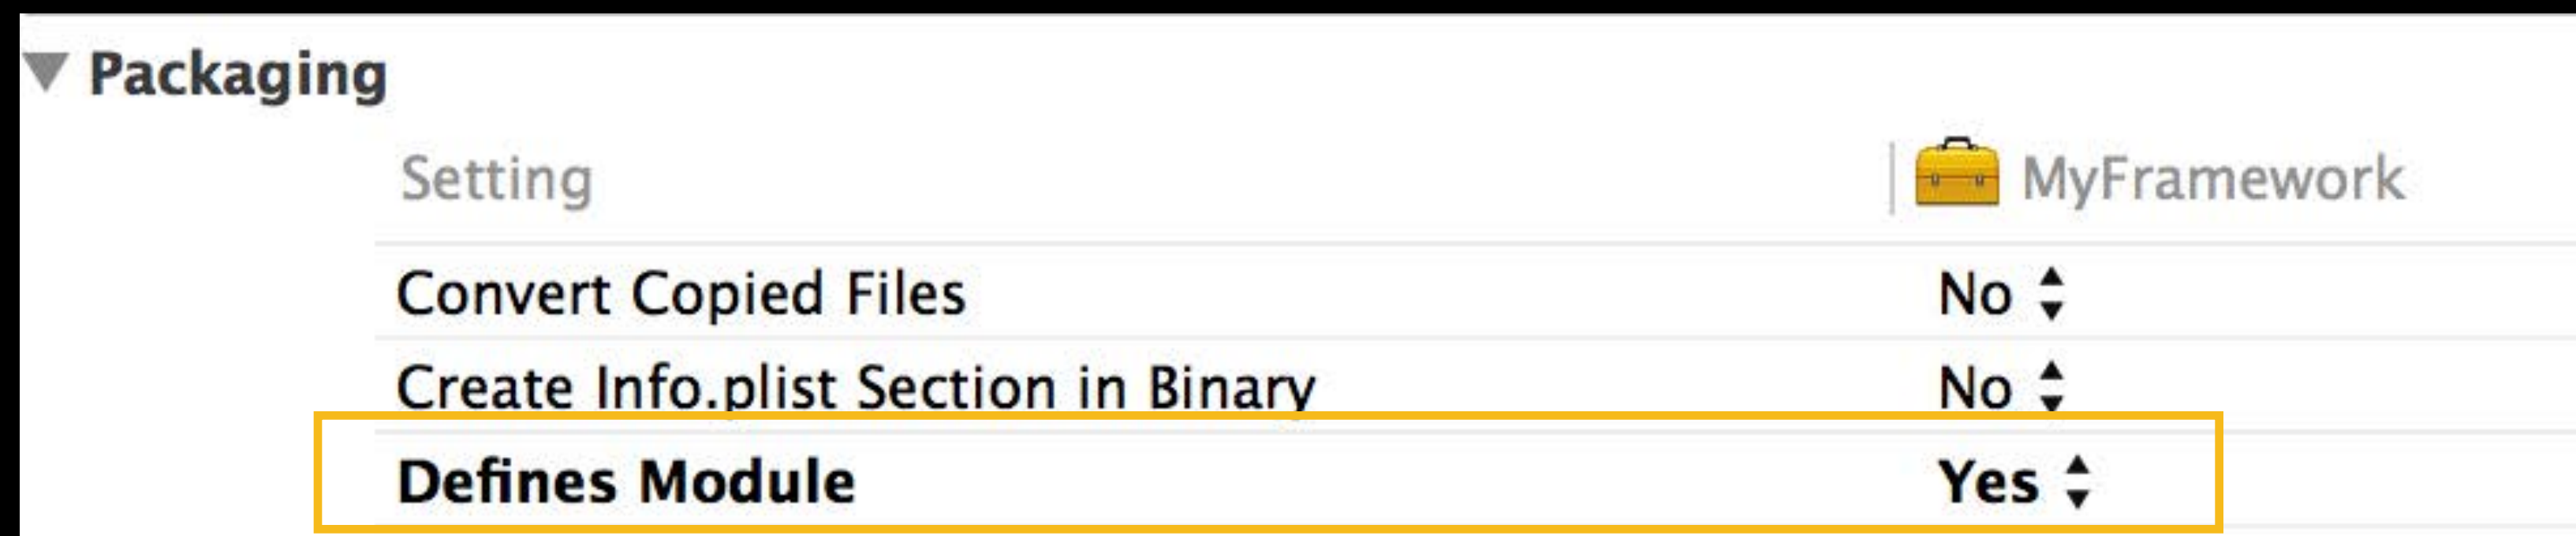


Now your C and Objective-C frameworks can define new modules

- Import your frameworks into Swift code
- Speed up compilation
- Avoid problems of textual inclusion

# Defining a Module

Provide an umbrella header that includes all the framework API



Custom module maps can describe more complex modules

For more information: <http://clang.llvm.org/docs/Modules.html>

# Importing a Module

```
@import MyFramework;           // import the module  
#import <MyFramework/MyFramework.h> // implicit modular import
```

## Guidelines

- Use @import when importing your framework in another target
- Use #import in the implementation to textually include the framework headers



# Module Rules

Do not expose nonmodular headers in the framework API

```
@import Cocoa;           // OK!
```

```
#import "Postgres.h"     // Only in the implementation, not the API
```

# Module Rules

Do not expose nonmodular headers in the framework API

```
@import Cocoa;           // OK!
```

```
#import "Posix.h" // Only in the implementation, not the API
```

# Module Rules

Do not expose nonmodular headers in the framework API

```
@import Cocoa;           // OK!  
#import "Resources.h"    // Only in the implementation, not the API
```

Modules can change the semantics

```
#define DEBUG 1  
@import MyFramework; // DEBUG ignored inside the framework
```

Use macro definitions on the command line (-DDEBUG=1) if necessary

# Summary: User-Defined Modules

Modules are the modern alternative to precompiled headers

- Fast compilation
- Clear semantics
- Swift interoperability

Define modules for your own frameworks!

# Profile Guided Optimization (PGO)

# Better Optimization via Profiling



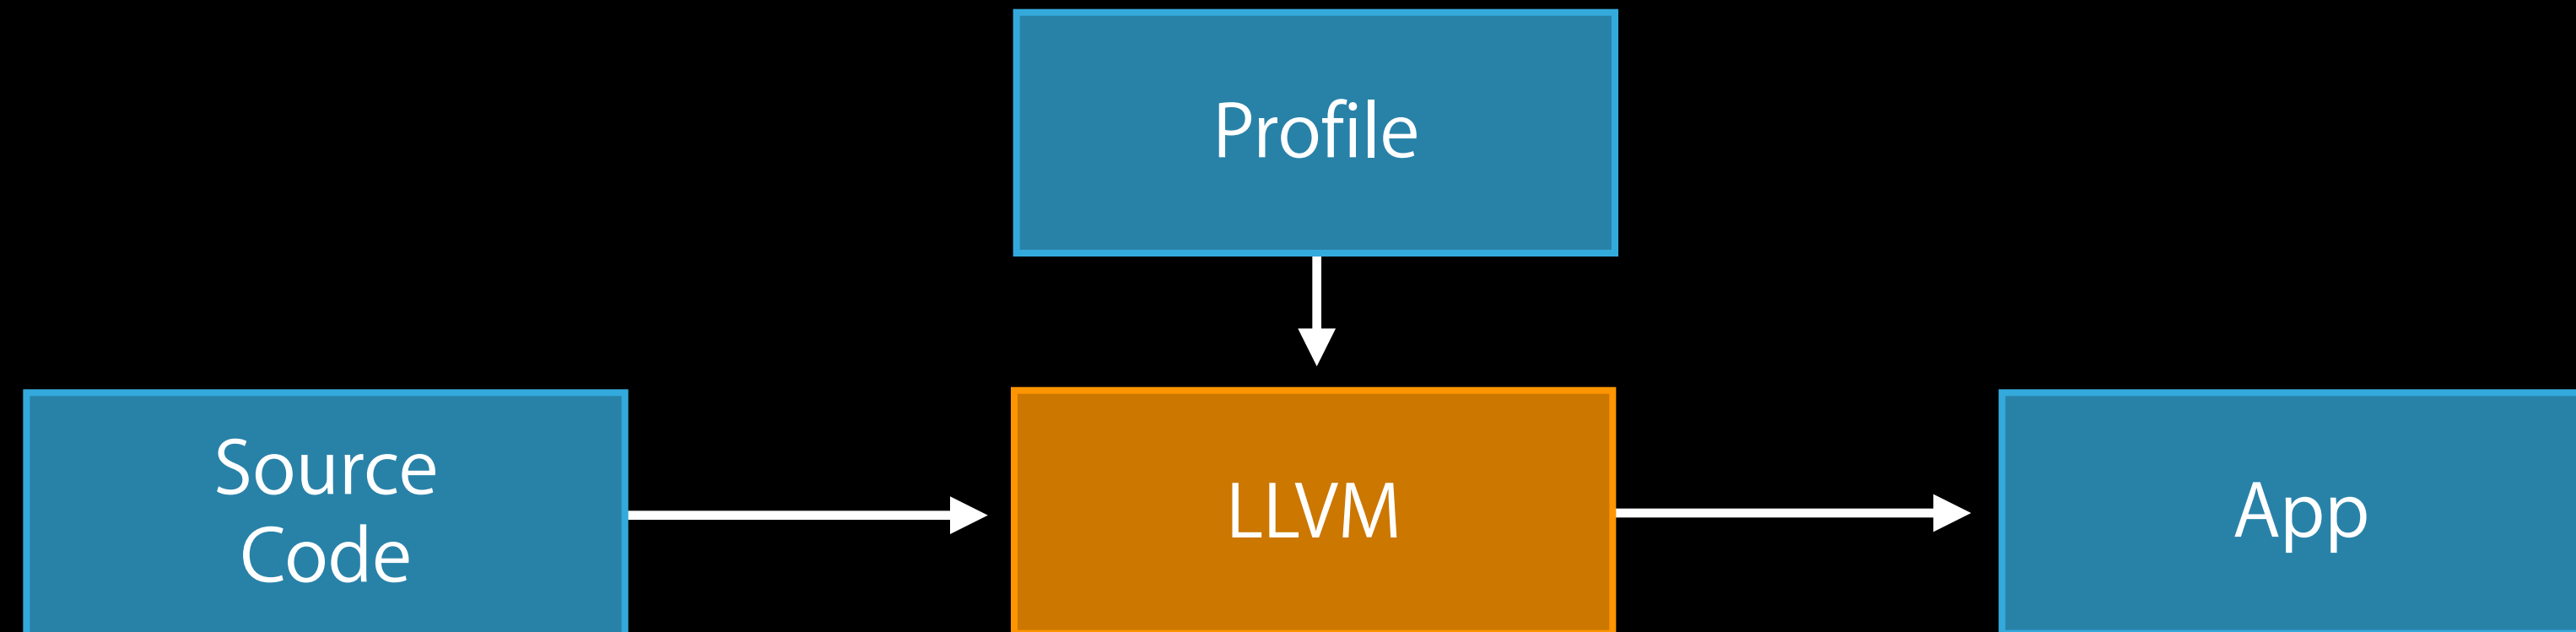
Problem: Compiler has to assume all code paths are equally likely



# Better Optimization via Profiling



Problem: Compiler has to assume all code paths are equally likely

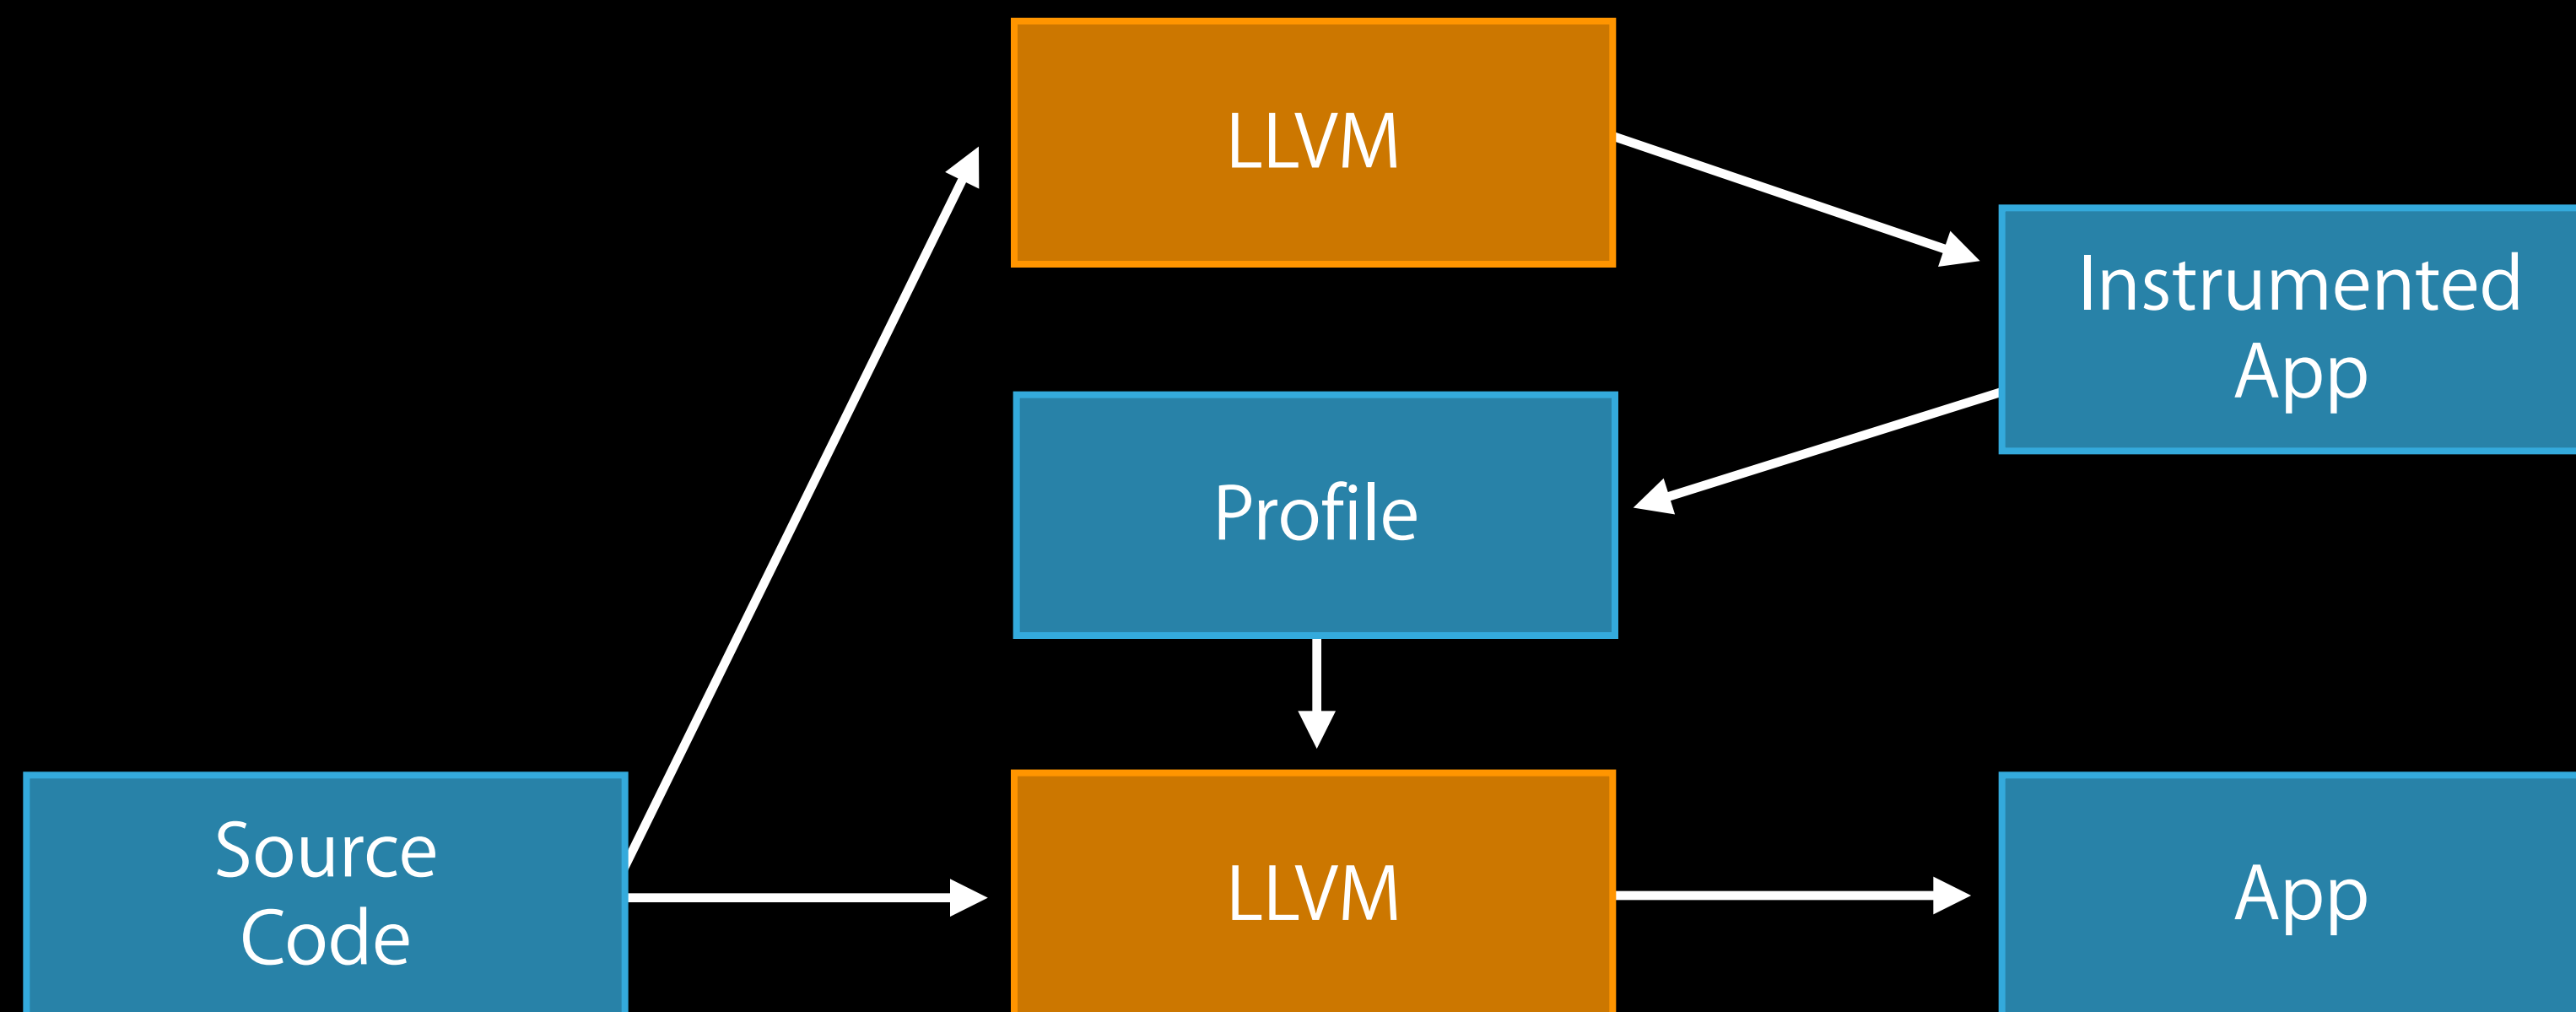


Solution: Use profile information to optimize for the common case

# Better Optimization via Profiling



Problem: Compiler has to assume all code paths are equally likely



Solution: Use profile information to optimize for the common case



# How Does Profiling Help?

Inline more aggressively for “hot” functions

Lay out most common code paths contiguously

Better register allocation

# Example: PGO Optimizations

```
for (ColoredObject *Obj : Objects)
    UpdatePosition(Obj);
```

```
void UpdatePosition(ColoredObject *Obj) {
    if (Obj->Color == Red) {
        // Red objects just move in a horizontal line.
        Obj->XCoord += 1;
    } else if (Obj->Color == Blue) {
        // Lots of complicated code here.
    }
}
```

# Example: PGO Optimizations

```
for (ColoredObject *Obj : Objects)
    UpdatePosition(Obj);
```

```
void UpdatePosition(ColoredObject *Obj) {
    if (Obj->Color == Red) {
        // Red objects just move in a horizontal line.
        Obj->XCoord += 1;
    } else if (Obj->Color == Blue) {
        // Lots of complicated code here.
    }
}
```

# Example: PGO Optimizations

```
for (ColoredObject *Obj : Objects)
    UpdatePosition(Obj);
```

```
void UpdatePosition(ColoredObject *Obj) {
    if (Obj->Color == Red) {
        // Red objects just move in a horizontal line.
        Obj->XCoord += 1;
    } else if (Obj->Color == Blue) {
        // Lots of complicated code here.
    }
}
```

# Example: PGO Optimizations

```
for (ColoredObject *Obj : Objects)
    UpdatePosition(Obj);
```

```
void UpdatePosition(ColoredObject *Obj) {
    if (Obj->Color == Red) {
        // Red objects just move in a horizontal line.
        Obj->XCoord += 1;
    } else if (Obj->Color == Blue) {
        // Lots of complicated code here.
    }
}
```

# Example: PGO Optimizations



Original code  
layout

# Example: PGO Optimizations



Original code  
layout

Inlining the  
“hot” function

# Example: PGO Optimizations



Original code  
layout



Inlining the  
"hot" function




# Example: PGO Optimizations



Original code  
layout



Inlining the  
“hot” function



Optimized code  
layout

# When Should You Use PGO?


Compiler already does a good job by default

PGO can provide even more improvements


- Requires some extra effort to collect a good profile
- Not all apps will benefit
- Worth the effort for very performance-sensitive code!

# PGO Performance

# PGO Performance



# PGO Performance



# How to Use PGO in Xcode

Step 1: Collect a profile

Step 2: Build with PGO

| ▼ Use Optimization Profile |   | <Multiple values> ⬆ |
|----------------------------|---|---------------------|
| Debug                      |   | No ⬆                |
| Release                    | + | Yes ⬆               |

# How to Use PGO in Xcode

Step 1: Collect a profile

Step 2: Build with PGO

| ▼ Use Optimization Profile |   | <Multiple values> ⬆ |
|----------------------------|---|---------------------|
| Debug                      |   | No ⬆                |
| Release                    | + | Yes ⬆               |

Compiler warns when the profile needs to be regenerated

⚠ Apple LLVM 6.0 Warning  
Profile data may be out of date: of 395 functions, 0 have no data and 1 has mismatched data that will be ignored

# Generating the Optimization Profile

Run Xcode's "Generate Optimization Profile" command

Xcode will build and run an instrumented version of your app

Important that you exercise all the performance-sensitive code



# Profiling with Tests

Can use performance tests to drive the profiling

- Reproducible results
- Requires good test coverage
- Tests should combine to reflect overall usage

Easy to evaluate the benefits of PGO

*Demo*

Profile Guided Optimization

# Summary: Profile Guided Optimization

Profile data enables better optimization

Be careful to collect good profiles and keep them updated

PGO can provide an extra boost in performance for many apps

# Vectorization in LLVM

Nadav Rotem

Manager, LLVM Performance Team

# Loop Vectorization Overview

Optimization for accelerating loops using vector instructions

```
for (int i = 0; i < 256; i++)  
    sum += A[i];
```



# Loop Vectorization Overview

Optimization for accelerating loops using vector instructions

```
for (int i = 0; i < 256; i++)  
    sum += A[i];
```



# Loop Vectorization Overview

Optimization for accelerating loops using vector instructions

```
for (int i = 0; i < 256; i++)  
    sum += A[i];
```



# Loop Vectorization Overview

Optimization for accelerating loops using vector instructions

```
for (int i = 0; i < 256; i++)  
    sum += A[i];
```





# What's New in the Loop Vectorizer

Better analysis of loops

PGO integration

Improved ARM64 and X86 code generation

Specialization of loop values


# Specialization of Loop Values

Most variables computed and known only at runtime

Vectorization often requires constant values

```
for (int i = 0; i < 256; i++)  
    sum += A[i * Step];
```

Nonconsecutive  
memory access



# Loop Specialization

Compiler creates multiple versions of the loop

Selects the correct version at runtime

```
if (Step == 1) {  
    for (i = 0; i < 256; ++i)  
        sum += A[i * 1];  
} else {  
    for (i = 0; i < 256; ++i)  
        sum += A[i * Step];  
}
```

← Check if Step == 1

← Vectorized loop

← Original loop

# SLP Vectorization



Superword Level Parallelism (SLP): Parallelism beyond loops

Combine multiple independent scalar calculations


Accelerates general code

# SLP Vectorization

```
struct Point {  
    double x, y;  
};  
  
void FtToCm(Point *P) {  
    P->x *= 30.48;  
    P->y *= 30.48;  
}
```

# SLP Vectorization

```
struct Point {  
    double x, y;  
};  
  
void FtToCm(Point *P) {  
    P->x *= 30.48;  
    P->y *= 30.48;  
}
```



Scalar


Load

Multiply

Store

# SLP Vectorization

```
struct Point {  
    double x, y;  
};  
  
void FtToCm(Point *P) {  
    P->x *= 30.48;  
    P->y *= 30.48;  
}
```



Scalar

Load

Multiply

Store

Load

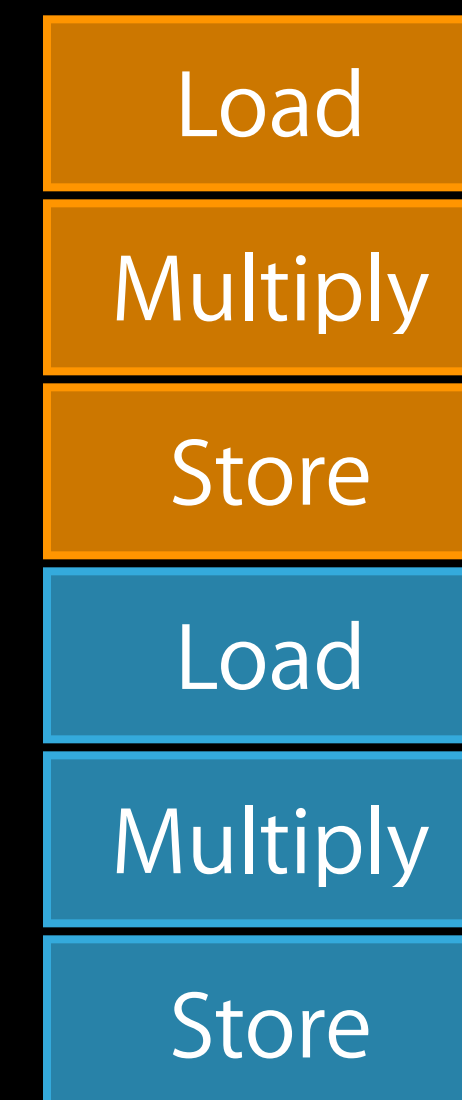
Multiply

Store


# SLP Vectorization

```
struct Point {  
    double x, y;  
};  
  
void FtToCm(Point *P) {  
    P->x *= 30.48;  
    P->y *= 30.48;  
}
```

Scalar




Vector





# SLP Vectorization Performance Gains

# SLP Vectorization Performance Gains



# Summary: Vectorization

Many improvements to loop vectorizer

New SLP vectorizer

Both vectorizers enabled for optimized builds


# Accelerating JavaScript Using LLVM

# WebKit Background

Uses an interpreter to execute JavaScript

Features two JITs (Just-in-Time compilers) to accelerate JavaScript

Tradeoff between quality of code and time spent in compiler



# JavaScript Is Evolving

Large compute intensive applications in JavaScript

Compile C++ to JavaScript and run in browsers

We need a better JavaScript compiler



+



# Fourth Tier LLVM (FTL)

LLVM as a fourth tier compiler

Compile functions that are executed many times

Takes longer to compile, but generates faster code





# Fourth Tier LLVM (FTL)

LLVM as a fourth tier compiler

Compile functions that are executed many times

Takes longer to compile, but generates faster code



# Compiling JavaScript Using LLVM

Compiling JavaScript is unlike C  
JavaScript is dynamically typed

```
function factorial(n) {  
    if (n === 0) {  
        return 1;  
    }  
    return n * factorial(n - 1);  
}
```

← Type of 'n'?

# Compiler Checks

Use information from previous runs to predict types

Inserts type checks and exception checks

If the checks fail, abort execution and return to interpreter

```
function factorial(n) {  
    if (n === 0) {  
        return 1;  
    }  
    return n * factorial(n - 1);  
}
```

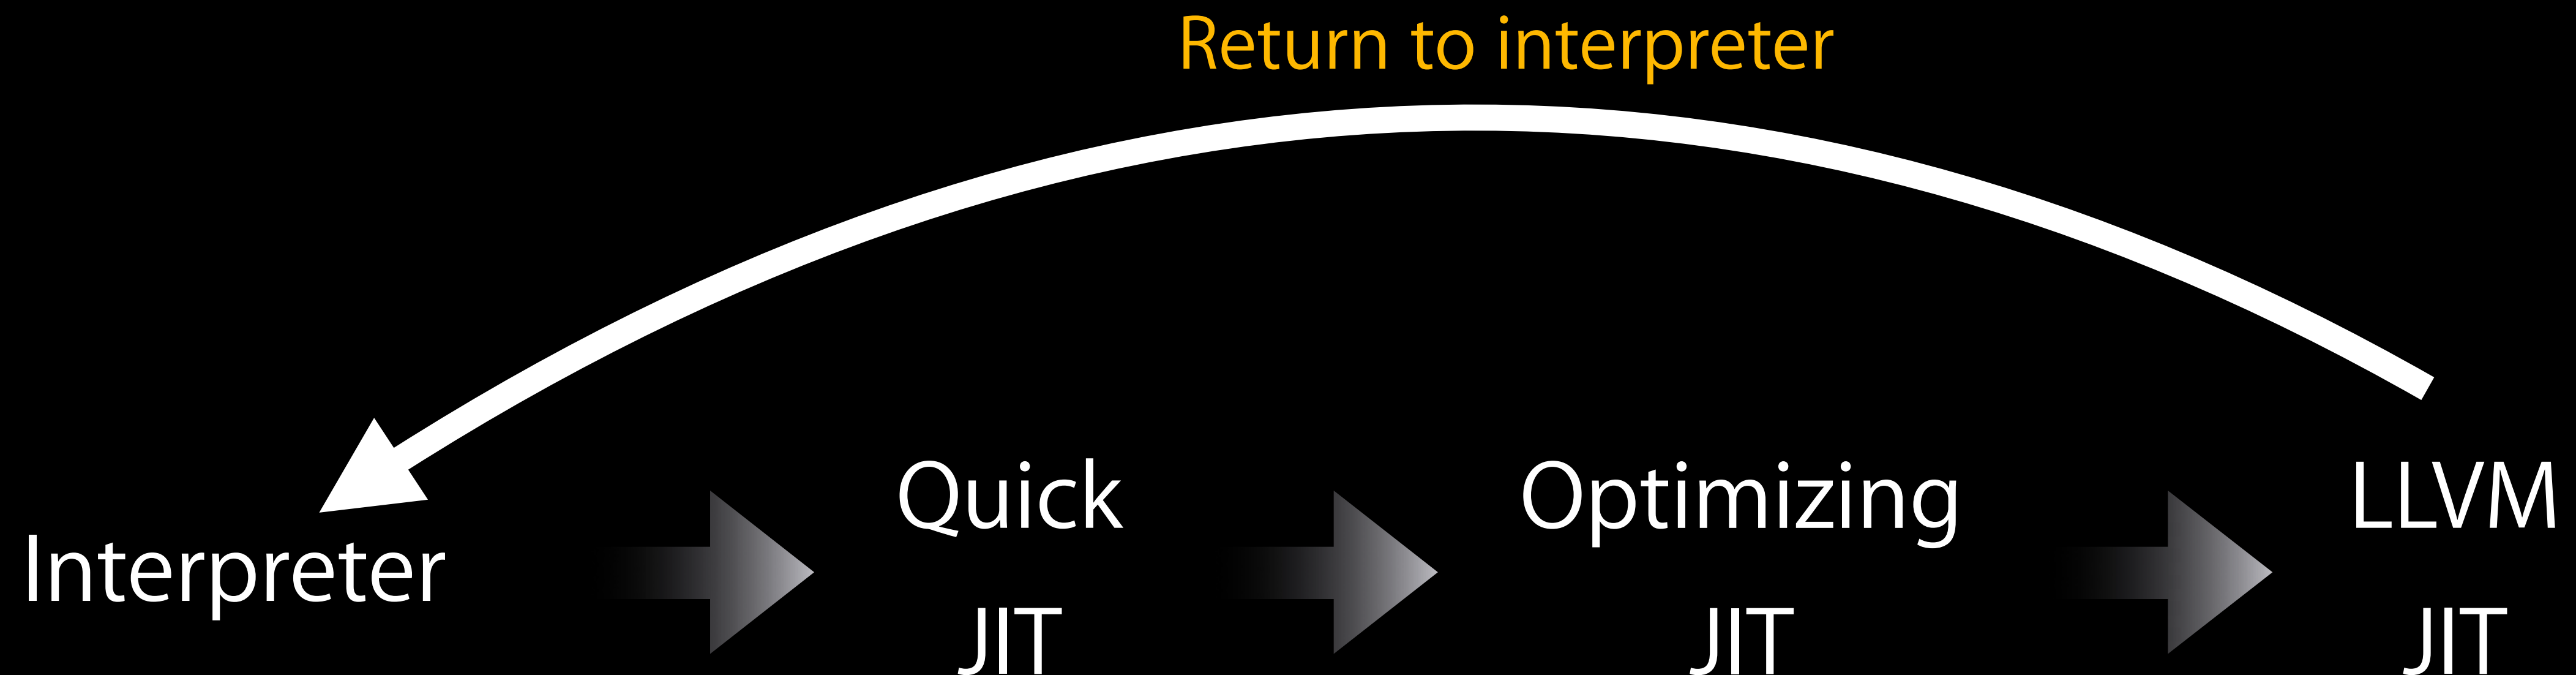
← 'n' is an Int

← Does not overflow

# Return to JavaScript Interpreter


Continue to interpret from the last valid checkpoint

“On-Stack Replacement” to migrate state to WebKit interpreter



# FTL Speedups

# FTL Speedups



# Summary: JavaScript FTL

New Fourth Tier LLVM (FTL) compiler for WebKit

Accelerates compute intensive JavaScript code

# Summary

LLVM helps modernize your code

- 64-bit iOS support
- Objective-C modernization tool
- User-defined modules

New LLVM optimizations increase performance

- Profile Guided Optimization
- Vectorizer advancements
- JavaScript Fourth Tier LLVM



# More Information

Dave DeLong  
Developer Tools Evangelist  
[delong@apple.com](mailto:delong@apple.com)

LLVM Project  
Open Source LLVM Project  
<http://llvm.org>

Apple Developer Forums  
<http://devforums.apple.com>

# Related Sessions

- |                                      |          |                  |
|--------------------------------------|----------|------------------|
| ● Integrating Swift with Objective-C | Presidio | Wednesday 9:00AM |
| ● Testing in Xcode 6                 | Marina   | Thursday 9:00AM  |

# Labs

|            |             |                  |
|------------|-------------|------------------|
| ● LLVM Lab | Tools Lab B | Wednesday 2:00PM |
| ● LLVM Lab | Tools Lab B | Thursday 2:00PM  |

